

In an AI-Native software organisation, agentic delegation is the default mode of production, not an accelerator applied to a human-authored baseline.

When production is cheap, the scarce resource is assurance.

	Flow & Recoverability	Verification	Context	Product Integrity	Control & Accountability
PILLARS → LEVELS ↓	A change reaches production, and can be undone, quickly and safely.	A change is known to be correct by someone who didn't write it.	The system is understandable to whoever — or whatever — works on it next.	What ships is deliberate, wanted, and coherent with what is already there.	Every change has a named human owner and a bounded blast radius.
3 Self-correcting The system improves its own hold. Evidence routinely changes the constraints. YOU MAY DELEGATE The safe-delegation boundary moves itself as evidence accumulates.	— Canary thresholds and rollout speed adjust from the observed failure rate for that class of change. — Fault injection runs continuously; a regression in restore time raises work before an incident does. — The path to production tightens or widens per change class without anyone running a project to do it.	— An escaped defect closes with the missing check added; the gate that failed is tightened as part of the fix rather than as follow-up work. — Flaky tests are quarantined and repaired automatically, so signal quality does not depend on human triage. — Check strictness and eval thresholds move with the measured risk of each change class.	— Context is regenerated from the system as it changes; documents failing their checks are corrected and re-proposed automatically. — Repeated agent failure in the same area raises a context fix rather than a retry. — Instructions are pruned on evidence of being unread, unused or contradictory, so the corpus doesn't bloat.	— Work that fails its hypothesis enters the retirement path on a timer unless someone argues to keep it — the default is removal, not survival. — The bar for entering the backlog moves with the measured hit rate. — Adding to the surface forces an explicit decision about what leaves.	— Permissions narrow themselves to observed need; unused grants expire without anyone revoking them. — Sandbox boundaries and approval thresholds adjust to each agent's demonstrated failure rate on that class of work. — Spend limits recalibrate against delivered value rather than being reset by argument.
2 Instrumented The system knows whether it's holding. The constraint is instrumented. YOU MAY DELEGATE Delegation calibrated per class of change from known failure rates.	— Time to restore is measured per class of change, not as a team-wide average. — The rollback mechanism is exercised on a schedule and its success rate is known. — The fraction of traffic exposed before a bad change is stopped is measured rather than assumed.	— Every escaped defect is traced to the specific check that should have caught it, and the gap is recorded. — Detection power is tracked over time, and the suite's false-positive rate is measured and budgeted — flakiness has a number and an owner. — First-pass acceptance of agent-authored changes is measured, and falling acceptance is treated as a signal about the constraints, not about the agent.	— The gap between system and documentation is measured — staleness, and time-to-first-useful-change for a newcomer or a fresh agent session. — Agent failures attributable to missing or wrong context are counted and traced to the specific document. — Duplication and coupling are trended; cognitive debt has a number that someone reviews.	— The proportion of shipped work that met its stated hypothesis is known, and quoted. — Every feature's usage is measured against its cost to keep. — The size of the product surface is a tracked number with a direction of travel.	— Permission grants are measured against use: what agents actually reached versus what they were permitted to. — Provenance completeness is measured, and unattributed changes are counted and chased. — Inference spend is measured per unit of delivered work, and blast radius is quantified per change class.
1 Constrained The system holds it. Enforced mechanically, whether or not anyone remembers. YOU MAY DELEGATE Agents work against machine-checkable constraints; humans review intent and sample the rest. Volume decouples from attention.	— The automated path is the only path to production; there is no manual route around it. — Changes ship decoupled from release — dark, behind flags, or progressively rolled out — so a bad change reaches a fraction of traffic before it is stopped. — Rollback is triggered by health signals rather than by a human noticing.	— Machine-checkable constraints — types, contracts, lint, security scans, performance budgets — run before a human sees the change, and agents iterate against them unattended. — The suite's detection power is evidenced by mutation score or deliberate fault injection, not inferred from coverage. — Behaviour that cannot be asserted deterministically is gated by evals with agreed thresholds.	— Documentation and instruction files are validated by the pipeline: stale commands, broken references and contradictory rules fail the build. — The context agents need is retrievable through a defined interface rather than scattered across tools. — Changes that duplicate existing code or widen coupling beyond a threshold are stopped by the system, not by a reviewer's memory.	— Work cannot enter the backlog without a stated hypothesis and a named owner; the system will not accept it otherwise. — Every launch carries a review date, and the outcome is recorded whether or not it flatters the decision. — Retirement is a funded routine with an owner; feature flags carry expiry dates and are removed on schedule.	— Agents are sandboxed by default; access to data, tools and networks is granted per task and expires. — Provenance — model, instruction version, accountable human — is recorded automatically and cannot be omitted. — Irreversible operations require a human decision by construction, and spend limits halt a runaway loop rather than surfacing on the bill.
0 Owned A person holds it. Individual diligence, applied reliably. YOU MAY DELEGATE Agents draft, humans approve everything. Attention is the throughput ceiling.	— Deployment is automated and any engineer on the team can run it unaided. — Rollback is a documented procedure with a named owner, deliberately exercised at least once. — Everything needed to rebuild the system is in version control, and environments are provisioned from it.	— Every change is reviewed by a named engineer who is accountable for it having merged. — The suite runs on every change and a red build blocks the merge. — Agent-authored changes are held to the same review standard as human-authored ones.	— Every component has a named owner and current documentation of its purpose and how to run it. — Agent instructions are maintained deliberately and reviewed like code. — Architectural decisions are recorded with the rationale that applied at the time.	— Every item of work has a named owner and a stated intended outcome, agreed before build starts. — Someone holds the authority to decline work, and exercises it. — The team can say what it declined last quarter, and why.	— Every change has a named human accountable for it, whether or not they wrote it. — Agents run under scoped credentials, with a documented list of what they may reach. — Agent authorship is attributed in the history; inference spend is attributed to teams and reviewed.
-1 Unowned Nothing holds the promise. Unowned and reactive. YOU MAY DELEGATE Delegate nothing. Agents amplify the mess.	— Deployment is a scheduled event, performed by hand, by whoever is available. — Rollback is a rebuild and redeploy of unknown duration; it has not been exercised this quarter. — Environments differ from production in ways nobody can enumerate.	— Failing tests are muted, skipped, or re-run until green. — Approvals arrive faster than the diff could plausibly have been read. — Nobody can say what the suite proves; coverage is quoted instead.	— Documentation exists and is known to be wrong; people ask a specific individual instead. — Agent instruction files have accumulated contradictory rules that nobody has pruned. — Changes are made by pattern-matching nearby code; no one can explain why the system is shaped as it is.	— Features ship with no stated purpose and no nameable owner. — Nothing has ever been retired; the product surface grows monotonically. — Requests become work items without anyone testing whether they are wanted.	— Agents run with the permissions of whoever launched them; nobody can enumerate what they can reach. — Agent-authored changes are indistinguishable from human ones after the fact. — Nobody knows what is being spent on inference, or by whom.